## THE RECORD LAYOUT OF A DATA SET IN SAS TRANSPORT (XPORT) FORMAT

## INTRODUCTION

```
All transport data set records are 80 bytes in length. If there is not
sufficient data to reach 80 bytes, then a record is padded with ASCII blanks
to 80 bytes. All character data are stored in ASCII, regardless of the
operating system. All integers are stored using IBM-style integer format,
and all floating-point numbers are stored using the IBM-style double
(truncated if the variable's length is less than 8). [An exception to this is
noted later.]
See the section "NUMERIC DATA FIELDS" for information on constructing
IBM-style doubles.
RECORD LAYOUT
1. The first header record consists of the following character string, in
    2. The first real header record uses the following layout:
       aaaaaaaabbbbbbbbcccccccddddddddeeeeeee
                                                                   ffffffffffffffff
In this record:
 -- aaaaaaaa and bbbbbbbb specify 'SAS
 -- ccccccc specifies 'SASLIB
  -- dddddddd specifies the version of the SAS(r) System under which the file was created.
 -- eeeeeeee specifies the operating system that creates the record.
 -- ffffffffffffffffffspecifies the date and time created, formatted as ddMMMyy:hh:mm:ss. Note
    that only a 2-digit year appears. If any program needs to read in this 2-digit year, be prepared
    to deal with dates in the 1900s or the 2000s.
Another way to consider this record is as a C structure:
    struct REAL_HEADER {
       char sas_symbol[2][8];
       char saslib[8];
       char sasver[8];
       char sas_os[8];
       char blanks[24];
       char sas_create[16];
       };
3. Second real header record
       ddMMMyy:hh:mm:ss
In this record, the string is the datetime modified. Most often, the datetime created and
datetime modified will always be the same. Pad with ASCII blanks to 80 bytes.
Note that only a 2-digit year appears. If any program needs to read in this
2-digit year, be prepared to deal with dates in the 1900s or the 2000s.
4. Member header records
  Both of these records occur for every member in the transport file.
       HEADER RECORD*******MEMBER HEADER RECORD!!!!!!!0000000000000000160000000140
       Note the 0140 that appears in the member header record above. This value specifies the size of
the variable descriptor (NAMESTR) record that is described later in this document. On the
VAX/VMS operating system, the value will be 0136 instead of 0140. This means that the
descriptor will be only 136 bytes instead of 140.
5. Member header data
                                                                    ffffffffffffffff
       aaaaaaaabbbbbbbbbbccccccccddddddddeeeeeeee
In this member header:
   - aaaaaaaa specifies 'SAS
 -- bbbbbbbb specifies the data set name
 -- ccccccc is SASDATA (if a SAS data set is being created)
 -- dddddddd specifies the version of the SAS System under which the file was created.
  -- eeeeeee specifies the operating system.
  -- fffffffffffffff is the datetime created, formatted as in previous headers.
Consider this C structure:
    struct REAL HEADER {
       char sas_symbol[8];
       char sas_dsname[8];
```

```
char sasdata[8];
       char sasver[8];
       char sas_osname[8];
       char blanks[24];
       char sas_create[16];
The second header record is as follows:
    ddMMMyy:hh:mm:ss
                                   In this record the datetime modified appears using the DATETIME16. format, followed by blanks
data set label and bbbbbbbb is the blank-padded data set type. Note that data set labels
can be up to 256 characters as of SAS 8, but only the first 40 characters are stored in the
second header record. Note also that only a 2-digit year appears in the datetime modified value.
If any program needs to read in this 2-digit year, be prepared to deal with dates in the 1900s
or the 2000s.
Consider the following C structure:
    struct SECOND HEADER {
       char dtmod_day[2]
       char dtmod month[3];
       char dtmod_year[2];
       char dtmod_colon1[1];
       char dtmod_hour[2];
       char dtmod colon2[1];
       char dtmod minute[2];
       char dtmod colon2[1];
       char dtmod_second[2];
       char padding[16];
       char dslabel[40];
       char dstype[8]
6. Namestr header record
One for each member.
    In this header record, xxxx is the number of variables in the data set, displayed with
blank-padded numeric characters. For example, for 2 variables, xxxx=0002. xxxx occurs
at offset 54 (base 0 as in C language use).
7. Namestr records
Each namestr field is 140 bytes long, but the fields are streamed together and
broken in 80-byte pieces. If the last byte of the last namestr field does not
fall in the last byte of the 80-byte record, the record is padded with ASCII
blanks to 80 bytes.
Here is the C structure definition for the namestr record:
  struct NAMESTR {
                              /* VARIABLE TYPE: 1=NUMERIC, 2=CHAR
   short ntype;
   short
           nhfun;
                              /* HASH OF NNAME (always 0)
                              /* LENGTH OF VARIABLE IN OBSERVATION
   short.
           nlng;
                              /* VARNUM
           nvar0;
   short
                              /* NAME OF VARIABLE
   char8
          nname;
                              /* LABEL OF VARIABLE
   char40 nlabel;
                              /* NAME OF FORMAT
/* FORMAT FIELD LENGTH OR 0
/* FORMAT NUMBER OF DECIMALS
   char8
           nform;
   short
           nfl;
           nfd;
   short
    short
           nfi;
                              /* 0=LEFT JUSTIFICATION, 1=RIGHT JUST
                              /* (UNUSED, FOR ALIGNMENT AND FUTURE)
/* NAME OF INPUT FORMAT
/* INFORMAT LENGTH ATTRIBUTE
           nfil1[2];
   char
   char8
           niform;
   short
           nifl;
                              /* INFORMAT NUMBER OF DECIMALS
   short
           nifd;
                              /* POSITION OF VALUE IN OBSERVATION
    long
           npos;
           rest[52];
                              /* remaining fields are irrelevant
   char
Note that the length given in the last 4 bytes of the member header record indicates the actual
number of bytes for the NAMESTR structure. The size of the structure listed above is 140 bytes. Under VAX/VMS, the size will be 136 bytes, meaning that the 'rest' variable may be truncated.
8. Observation header
    HEADER RECORD******OBS
                               9. Data records
Data records are streamed in the same way that namestrs are. There is ASCII blank padding at
```

```
the end of the last record if necessary. There is no special trailing record.
MISSING VALUES
Missing values are written out with the first byte (the exponent) indicating
the proper missing values. All subsequent bytes are 0x00. The first byte is:
               bvte
     type
               0x5f
      . A
               0 \times 41
      .в
               0x42
               0x5a
A SAMPLE SESSION TO SHOW A TRANSPORT DATA SET
This session was run on a ASCII-based system to demonstrate the creation
and record layout of a transport data set.
     $ sas606
      1? libname xxx sasv5xpt 'xxx.dat';
    NOTE: Libref XXX was successfully assigned as follows:
                          XPORT
           Engine:
           Physical Name: xxx.dat
      8? data temp; input x y $ @@; cards;
      9> 1 a 2 B . . .a *
     10> run;
     NOTE: SAS went to a new line when INPUT statement reached past the end of a
           line.
     NOTE: The data set WORK.TEMP has 4 observations and 2 variables.
    NOTE: The DATA statement used 10 seconds.

NOTE: The DATA statement used 1 seconds cpu time.
      11? data temp; set temp;
     12? format x date7.; label y='character variable'; run;
     NOTE: The data set WORK.TEMP has 4 observations and 2 variables.
    NOTE: The DATA statement used 12 seconds.
     NOTE: The DATA statement used 2 seconds cpu time.
     13? proc print data=temp; format x y ; run;
                                     The SAS System 10:17 Thursday, April 13, 1989
                                     OBS
                                          X Y
                                       1
                                       2
                                            2
                                                  В
                                       4
     NOTE: The PROCEDURE PRINT used 3 seconds.
     NOTE: The PROCEDURE PRINT used 1 seconds cpu time.
      14? proc print data=temp; run;
                                     The SAS System 10:17 Thursday, April 13, 1989
                                     OBS
                                      1
                                            02JAN60
                                            03JAN60
                                                        В
        NOTE: The PROCEDURE PRINT used 2 seconds.
        NOTE: The PROCEDURE PRINT used less than 1 second cpu time.
        15? proc contents; run;
                                        The SAS System 10:17 Thursday, April 13, 1989
                                   CONTENTS PROCEDURE
           Data Set Name: WORK.TEMP
                                                     Observations:
           Member Type: DATA
                                                     Variables:
           Engine:
                          V606
                                                     Indexes:
           Created:
                          13APR89:10:19:15
                                                    Observation Length: 16
           Last Modified: 13APR89:10:19:15
                                                    Deleted Observations: 0
           Data Set Type:
                                                    Compressed:
           Label:
                 -----Alphabetic List of Variables and Attributes-----
```

```
Type
        Variable
                   Len
                        Pos
                            Format
                                 Label
                          DATE7.
               Num
               Char
                                  character variable
             ----Engine/Host Dependent Information----
                     The SAS System 10:17 Thursday, April 13, 1989
                     CONTENTS PROCEDURE
                 Data Set Page Size:
                                4096
                 Number of Data Set Pages: 1
                 First Data Page:
                 Max Obs per Page:
                 Obs in First Data Page:
                 FILETYPE:
                                REGULAR
    NOTE: The PROCEDURE CONTENTS used 7 seconds.
    NOTE: The PROCEDURE CONTENTS used 2 seconds cpu time.
     16? data xxx.abc; set; run;
    NOTE: The data set XXX.ABC has 4 observations and 2 variables.
    NOTE: The DATA statement used 2 seconds.
    NOTE: The DATA statement used 1 seconds cpu time.
     20? options ls=132;
     21? data _null_; infile 'xxx.dat' recfm=f lrecl=80; input x $char80.;list;run;
  RUILE:---+---1----2---+---3----+---4---+---5---+---6---+---7---+---8----+
     SASLIB 6.06
     SAS
          SAS
                        bsd4.2
                                           13APR89:10:20:06
     SAS
          ABC
              SASDATA 6.06
                       bsd4.2
                                           13APR89:10:20:06
     13APR89:10:20:06
     CHAR .
  CHAR
  HEADER RECORD*****OBS
                     A .....B
  CHAR A....a
                                 A....*
  NOTE: The infile 'xxx.dat' is:
      FILENAME=//HOBBITT/UDR/LANGSTON/COM/XXX.DAT
  NOTE: 14 records were read from the infile 'xxx.dat'.
  NOTE: The DATA statement used 4 seconds.
  NOTE: The DATA statement used less than 1 second cpu time.
  NOTE: SAS Institute Inc., SAS Circle, PO Box 8000, Cary, NC 27512-8000
In case you're not familiar with the LIST output from the DATA step, here's
a brief explanation:
If the record has any unprintable characters, LIST output generates three lines of
1) the record itself, printing everything that's printable and using dots for everything else
2) the upper nibble of each byte in hex
3) the lower nibble of each byte in hex.
Consider, then, record 9, which has some printable and unprintable characters:
  The first 8 bytes are unprintable, since dots appear. Those first 8 bytes, reading
in sequential hex format, would be 00010000 00080001. The next 56 bytes are
printable. Then, we have 00070000 00000000. The remaining 8 bytes are ASCII blanks.
NUMERIC DATA FIELDS
All numeric data fields in the transport file are stored as floating-point numbers.
```

output:

All floating-point numbers in the file are stored using the IBM mainframe representation. If your application is to read from or write to transport files, it will be necessary to convert native floating-point numbers to or from the transport representation.

Most platforms use the IEEE representation for floating-point numbers. Some of these platforms store the floating-point numbers in reversed byte order from other platforms. For the sake of nomenclature, we will call these platforms "big endian" and "little endian" platforms.

A big-endian environment stores integers with the lowest-significant byte at a higher address in memory. Likewise, an IEEE platform is big endian if the first byte of the exponent is stored in a lower address than the first byte of the mantissa. For example, the HP series machines store a floating-point 1 as 3F F0 00 00 00 00 00 (the bytes in hex), while an IBM PC stores a 1 as 00 00 00 00 00 00 F0 3F. The bytes are the same, just reversed. Therefore, the HP is considered big endian and the PC is considered little endian.

This is a partial list of the categories of machines on which the SAS System runs:

| Hardware      | Operating<br>Systems | Float Type | Endian |
|---------------|----------------------|------------|--------|
| IBM mainframe | MVS,CMS,VSE          | IBM        | big    |
| DEC Alpha     | AXP/VMS,DEC UNIX     | IEEE       | little |
| HP            | HP-UX                | IEEE       | big    |
| Sun           | Solaris I,II         | IEEE       | big    |
| RS/6000       | AIX                  | IEEE       | big    |
| IBM PC        | Windows,OS/2,IABI    | IEEE       | little |
|               |                      |            |        |

Not included is VAX, which uses a different floating-point representation than either IBM mainframe or IEEE.

## PROVIDED SUBROUTINES

To assist you in reading and/or writing transport files, we are providing routines to convert from IEEE representation (either big endian or little endian) to transport representation and back again. The source code for these routines is provided at the end of this document. Note that the source code is provided as is, and as a convenience to those needing to read and/or write transport files. The source code has been tested on HP-UX, DEC UNIX, IBM PC, and MVS.

The routine to use is cnxptiee. This routine will convert in either direction, either to or from transport. Its usage is as follows:

```
/* rc = cnxptiee(from,fromtype,to,totype);
     from
                        pointer to a floating point value
                        type of floating point value (see below) pointer to target area
     fromtype
     to
                        type of target value (see below)
     totype
/* floating point types:
     0
                        native floating point
                        IBM mainframe (transport representation)
     1
                        Big endian IEEE floating point
                        Little endian IEEE floating point
/* rc = cnxptiee(from,0,to,1); native -> transport
/* rc = cnxptiee(from,0,to,2); native -> Big endian IEEE
/* rc = cnxptiee(from,0,to,3); native -> Little endian IEEE
/* rc = cnxptiee(from,1,to,0); transport -> native
/* rc = cnxptiee(from,1,to,2); transport -> Big endian IEEE
                                 transport -> Little endian IEEE
/* rc = cnxptiee(from,1,to,3);
/* rc = cnxptiee(from,2,to,0); Big endian IEEE -> native
/* rc = cnxptiee(from,2,to,1); Big endian IEEE -> transport
/* rc = cnxptiee(from,2,to,3); Big endian IEEE -> Little endian IEEE
/* rc = cnxptiee(from, 3, to, 0); Little endian IEEE -> native
/* rc = cnxptiee(from,3,to,1); Little endian IEEE -> transport
/* rc = cnxptiee(from,3,to,2); Little endian IEEE -> Big endian IEEE */
```

The "native" representation is whatever is appropriate for the host machine. Most likely you will be using that mode.

The testieee.c routine is supplied here to demonstrate how the cnxptiee is used. It is also useful to ensure that the cnxptiee routine works in your environment.

Note that there are several symbols that can be defined when compiling the ieee.c file. These symbols are FLOATREP, BIG\_ENDIAN, and LITTLE\_ENDIAN.

FLOATREP should be set to one of the following strings:

```
-- CN_TYPE_IEEEB
                          Big endian IEEE
                         Little endian IEEE
 -- CN TYPE TEEEL
  -- CN TYPE XPORT
                          Transport format (for example, IBM)
If BIG_ENDIAN is defined, it is assumed that the platform is big endian. If LITTLE_ENDIAN is
defined, it is assumed that the platform is little endian. Do not define both of them.
If FLOATREP is not defined, the proper value is determined at run time. Although this works,
it incurs additional overhead that can increase CPU time with large files. Use the FLOATREP
symbol to improve efficiency. Likewise, if neither BIG_ENDIAN nor LITTLE_ENDIAN is defined, the
proper orientation is determined at run time. It is much more efficient to supply the proper
definition at compile time.
As an example, consider this command on HP-UX:
     cc testieee.c ieee.c -DFLOATREP=CN TYPE IEEEB -DBIG ENDIAN
Also consider the corresponding command on DEC UNIX:
     cc testieee.c ieee.c -DFLOATREP=CN TYPE IEEEL -DLITTLE ENDIAN
Here is the correct output from the testieee run:
     Native -> Big endian IEEE match count = 4 (should be 4).
     Native -> Little endian IEEE match count = 4 (should be 4).
     Native -> Transport match count = 4 (should be 4).
     Transport -> Big endian IEEE match count = 4 (should be 4).
     Transport -> Little endian IEEE match count = 4 (should be 4).
     Transport -> Native match count = 4 (should be 4).
Big endian IEEE -> Little endian IEEE match count = 4 (should be 4).
     Big endian IEEE -> Transport match count = 4 (should be 4). Big endian IEEE -> Native match count = 4 (should be 4).
     Little endian IEEE -> Big endian IEEE match count = 4 (should be 4).
     Little endian IEEE -> Transport match count = 4 (should be 4).
     Little endian IEEE -> Native match count = 4 (should be 4).
Here is the source code for the test program, testieee.c
       -----testieee.c------
     #define CN_TYPE_NATIVE 0
     #define CN_TYPE_XPORT 1
        #define CN_TYPE_IEEEB 2
        #define CN_TYPE_IEEEL 3
     void tohex();
     #define N TESTVALS 4
     static char xpt_testvals[N_TESTVALS][8] = {
        {0x41,0x10,0x00,0x00,0x00,0x00,0x00,0x00}, /* 1 */
{0xc1,0x10,0x00,0x00,0x00,0x00,0x00,0x00}, /* -1 */
        static char ieeeb_testvals[N_TESTVALS][8] = {
        static char ieeel_testvals[N_TESTVALS][8] = {
        {0x00,0x00,0x00,0x00,0x00,0x00,0xf0,0x3f}, /* 1 */
        static double native[N_TESTVALS] =
        {1,-1,0,2};
     #define N_MISSINGVALS 3
        static char missingvals[N_MISSINGVALS][8] = {
    {0x2e,0x00,0x00,0x00,0x00,0x00,0x00,0x00}, /* std missing */
    {0x41,0x00,0x00,0x00,0x00,0x00,0x00,0x00}, /* .A */
    {0x5A,0x00,0x00,0x00,0x00,0x00,0x00,0x00} /* .Z */
     /* rc = cnxptiee(from,0,to,1); native -> transport */
        /* rc = cnxptiee(from,0,to,2); native -> Big endian IEEE */
        /* rc = cnxptiee(from,0,to,3); native -> Little endian IEEE */
        /* rc = cnxptiee(from,1,to,0); transport -> native */
        /* rc = cnxptiee(from,1,to,2); transport -> Big endian IEEE */
        /* rc = cnxptiee(from,1,to,3); transport -> Little endian IEEE */
```

```
/* rc = cnxptiee(from,2,to,0); Big endian IEEE -> native */
/* rc = cnxptiee(from,2,to,1); Big endian IEEE -> transport */
/* rc = cnxptiee(from,2,to,3); Big endian IEEE -> Little endian IEEE */
   /* rc = cnxptiee(from,3,to,0); Little endian IEEE -> native */
   /* rc = cnxptiee(from,3,to,1); Little endian IEEE -> transport */
   /* rc = cnxptiee(from,3,to,2); Little endian IEEE -> Big endian IEEE */
main()
   char to[8];
   int i, matched;
   char hexdigits[17];
for (i=matched=0;i<N_TESTVALS;i++) {
   cnxptiee(&native[i],CN_TYPE_NATIVE,to,CN_TYPE_IEEEB);
   matched += (memcmp(to,ieeeb_testvals[i],8) == 0);
   printf("Native -> Big endian IEEE match count = %d (should be %d).\n",
   matched, N_TESTVALS);
for (i=matched=0;i<N TESTVALS;i++)
   cnxptiee(&native[i],CN_TYPE_NATIVE,to,CN_TYPE_IEEEL);
   matched += (memcmp(to,ieeel_testvals[i],8) == 0);
   printf("Native -> Little endian IEEE match count = %d (should be %d).\n",
   matched,N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {
   cnxptiee(&native[i],CN_TYPE_NATIVE,to,CN_TYPE_XPORT);
   matched += (memcmp(to,xpt_testvals[i],8) == 0);
   printf("Native -> Transport match count = %d (should be %d).\n",
   matched, N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {</pre>
   cnxptiee(xpt_testvals[i],CN_TYPE_XPORT,to,CN_TYPE_IEEEB);
   matched += (memcmp(to,ieeeb testvals[i],8) == 0);
   printf("Transport -> Big endian IEEE match count = %d (should be %d).\n",
   matched, N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {
   cnxptiee(xpt_testvals[i],CN_TYPE_XPORT,to,CN_TYPE_IEEEL);
   matched += (memcmp(to,ieeel_testvals[i],8) == 0);
   printf("Transport -> Little endian IEEE match count = %d \
   (should be %d).\n"
   matched, N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {
   cnxptiee(xpt_testvals[i],CN_TYPE_XPORT,to,CN_TYPE_NATIVE);
   matched += (memcmp(to,&native[i],8) == 0);
   printf("Transport -> Native match count = %d (should be %d).\n",
   matched, N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {
   cnxptiee(ieeeb_testvals[i],CN_TYPE_IEEEB,to,CN_TYPE_IEEEL);
   matched += (memcmp(to,ieeel_testvals[i],8) == 0);
   printf("Big endian IEEE -> Little endian IEEE match count = %d \ (should be %d).\n",
   matched,N_TESTVALS);
for (i=matched=0;i<N TESTVALS;i++) {
   cnxptiee(ieeeb_testvals[i],CN_TYPE_IEEEB,to,CN_TYPE_XPORT);
   matched += (memcmp(to,xpt_testvals[i],8) == 0);
   printf("Big endian IEEE -> Transport match count = %d (should be %d).\n",
   matched, N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {
   cnxptiee(ieeeb_testvals[i],CN_TYPE_IEEEB,to,CN_TYPE_NATIVE);
   matched += (memcmp(to,&native[i],8) == 0);
   printf("Big endian IEEE -> Native match count = %d (should be %d).\n",
   matched, N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {
   cnxptiee(ieeel testvals[i],CN TYPE IEEEL,to,CN TYPE IEEEB);
   matched += (memcmp(to,ieeeb_testvals[i],8) == 0);
   printf("Little endian IEEE -> Big endian IEEE match count = %d \setminus
   (should be %d).\n"
   matched, N_TESTVALS);
```

```
for (i=matched=0;i<N_TESTVALS;i++) {</pre>
   cnxptiee(ieeel_testvals[i],CN_TYPE_IEEEL,to,CN_TYPE_XPORT);
   matched += (memcmp(to,xpt_testvals[i],8) == 0);
   printf("Little endian IEEE -> Transport match count = %d (should be %d).\n",
   matched, N_TESTVALS);
for (i=matched=0;i<N_TESTVALS;i++) {</pre>
   cnxptiee(ieeel_testvals[i],CN_TYPE_IEEEL,to,CN_TYPE_NATIVE);
   matched += (memcmp(to,&native[i],8) == 0);
   printf("Little endian IEEE -> Native match count = %d (should be %d).\n",
   matched, N_TESTVALS);
void tohex(bytes,hexchars,length)
   unsigned char *bytes;
   char *hexchars;
   int length;
   static char *hexdigits = "0123456789ABCDEF";
   int i;
   for (i=0;i<length;i++) {
 *hexchars++ = hexdigits[*bytes >> 4];
 *hexchars++ = hexdigits[*bytes++ & 0x0f];
    *hexchars = 0;
-----ieee.c-----ieee.c-----
   #define CN_TYPE_NATIVE 0
   #define CN_TYPE_XPORT 1
    #define CN_TYPE_IEEEB 2
   #define CN_TYPE_IEEEL 3
int cnxptiee();
    void xpt2ieee();
   void ieee2xpt();
#ifndef FLOATREP
    #define FLOATREP get_native()
   int get_native();
   #endif
   /* rc = cnxptiee(from,fromtype,to,totype); */
   /* where */
   /* from pointer to a floating point value */
   /* fromtype type of floating point value (see below) */
    /* to pointer to target area */
    /* totype type of target value (see below) */
   /* floating point types: */
/* 0 native floating point */
   /* 1 IBM mainframe (transport representation) */
    /* 2 Big endian IEEE floating point */
   /* 3 Little endian IEEE floating point */
   /* rc = cnxptiee(from,0,to,1); native -> transport */
   /* rc = cnxptiee(from,0,to,2); native -> Big endian IEEE */
/* rc = cnxptiee(from,0,to,3); native -> Little endian IEEE */
    /* rc = cnxptiee(from,1,to,0); transport -> native */
   /* rc = cnxptiee(from,1,to,0); transport -> hative "/

/* rc = cnxptiee(from,1,to,2); transport -> Big endian IEEE */

/* rc = cnxptiee(from,1,to,3); transport -> Little endian IEEE */

/* rc = cnxptiee(from,2,to,0); Big endian IEEE -> native */

/* rc = cnxptiee(from,2,to,1); Big endian IEEE -> transport */

/* rc = cnxptiee(from,2,to,3); Big endian IEEE -> Little endian IEEE */
    /* rc = cnxptiee(from,3,to,0); Little endian IEEE -> native */
    /* rc = cnxptiee(from,3,to,1); Little endian IEEE -> transport */
    /* rc = cnxptiee(from,3,to,2); Little endian IEEE -> Big endian IEEE */
int cnxptiee(from,fromtype,to,totype)
   char *from;
int fromtype;
   char *to;
   int totype;
   char temp[8];
   int i;
if (fromtype == CN_TYPE_NATIVE) {
   fromtype = FLOATREP;
```

```
switch(fromtype) {
   case CN_TYPE_IEEEL :
   if (totype == CN_TYPE_IEEEL)
   for (i=7;i>=0;i--) {
temp[7-i] = from[i];
   from = temp;
   fromtype = CN_TYPE_IEEEB;
   /* break intentionally omitted */
   case CN_TYPE_IEEEB :
   /* break intentionally omitted */
   case CN_TYPE_XPORT :
   break;
   default:
   return(-1);
   if (totype == CN_TYPE_NATIVE) {
totype = FLOATREP;
   switch(totype) {
   case CN_TYPE_XPORT :
   case CN_TYPE_IEEEB :
   case CN_TYPE_IEEEL :
   break;
   default:
   return(-2);
   if (fromtype == totype) {
  memcpy(to,from,8);
return(0);
   switch(fromtype) {
   case CN_TYPE_IEEEB :
   if (totype == CN_TYPE_XPORT)
   ieee2xpt(from,to);
   else memcpy(to,from,8);
   break;
   case CN_TYPE_XPORT :
   xpt2ieee(from,to);
   break;
   if (totype == CN_TYPE_IEEEL) {
   memcpy(temp,to,8);
   for (i=7;i>=0;i--) {
   to[7-i] = temp[i];
   return(0);
int get_native() {
static char float_reps[][8] = {
   static double one = 1.00;
int i,j;
  i ; j = sizeof(float_reps)/8;
for (i=0;i<j;i++) {
  if (memcmp(&one,float_reps+i,8) == 0)
  return(i+1);</pre>
   return(-1);
#ifdef BIG_ENDIAN
   #define REVERSE(a,b)
   #endif
#ifdef LITTLE_ENDIAN
   #define DEFINE_REVERSE
   void REVERSE();
   #endif
#if !defined(DEFINE_REVERSE) && !defined(REVERSE)
   #define DEFINE_REVERSE
   void REVERSE();
   #endif
void xpt2ieee(xport,ieee)
```

```
unsigned char *xport;
   unsigned char *ieee;
   char temp[8];
   register int shift;
register int nib;
   unsigned long ieeel, ieee2;
   unsigned long xport1 = 0;
   unsigned long xport2 = 0;
memcpy(temp,xport,8);
   memset(ieee,0,8);
if (*temp && memcmp(temp+1,ieee,7) == 0) {
  ieee[0] = ieee[1] = 0xff;
   ieee[2] = \sim(*temp);
   return;
memcpy(((char *)&xport1)+sizeof(unsigned long)-4,temp,4);
   REVERSE(&xport1,sizeof(unsigned long));
   memcpy(((char *)&xport2)+sizeof(unsigned long)-4,temp+4,4);
   REVERSE(&xport2,sizeof(unsigned long));
 /* Translate IBM format floating point numbers into IEEE */
   /* format floating point numbers. */
   /* IEEE format: */
   /* 6 5 0 */
   /* 3 1 0 */
   /* */
   /* SEEEEEEEEEEMMMM ..... MMMM */
   /* Sign bit, 11 bits exponent, 52 bit fraction. Exponent is ^{*}/
   /* excess 1023. The fraction is multiplied by a power of 2 of */
   /* the actual exponent. Normalized floating point numbers are */
   /* represented with the binary point immediately to the left */
/* of the fraction with an implied "1" to the left of the */
   /* binary point. */
   /* IBM format: */
   /* 6 5 0 */
   /* 3 1 0 */
   /* SEEEEEEEMMMM ..... MMMM */
   /* Sign bit, 7 bit exponent, 56 bit fraction. Exponent is */
   /* excess 64. The fraction is multiplied by a power of 16 of */
   /* the actual exponent. Normalized floating point numbers are */
/* represented with the radix point immediately to the left of */
   /* the high order hex fraction digit. */
   /* How do you translate from IBM format to IEEE? */
   /* Translating back to ieee format from ibm is easier than */
   /* going the other way. You lose at most, 3 bits of fraction, */
/* but nothing can be done about that. The only tricky parts */
   /* are setting up the correct binary exponent from the ibm */
   /* hex exponent, and removing the implicit "1" bit of the ieee*/
   /* fraction (see vzctdbl). We must shift down the high order */
   /* nibble of the ibm fraction until it is 1. This is the */
/* implicit 1. The bit is then cleared and the exponent */
/* adjusted by the number of positions shifted. A more */
   /* thorough discussion is in vzctdbl.c. */
 /st Get the first half of the ibm number without the exponent st/
   /* into the ieee number *
   ieeel = xport1 & 0x00ffffff;
 /* get the second half of the ibm number into the second half */
      of the ieee number . If both halves were 0. then just */
   /* return since the ieee number is zero. */
   if ((!(ieee2 = xport2)) && !xport1)
   return;
 /* The fraction bit to the left of the binary point in the */
   /* ieee format was set and the number was shifted 0, 1, 2, or */ /* 3 places. This will tell us how to adjust the ibm exponent */
   /\,^{\star} to be a power of 2 ieee exponent and how to shift the ^{\star}/\,
   /* fraction bits to restore the correct magnitude. */
 if ((nib = (int)xport1) & 0x00800000)
```

```
shift = 3;
   else
   if (nib & 0x00400000)
   shift = 2;
   if (nib & 0 \times 00200000)
   shift = 1;
   else
   shift = 0;
 if (shift)
   {
/* shift the ieee number down the correct number of places*/
   /* then set the second half of the ieee number to be the */
   /* second half of the ibm number shifted appropriately, */
   /st ored with the bits from the first half that would have st/
   /* been shifted in if we could shift a double. All we are */
/* worried about are the low order 3 bits of the first */
   /* half since we're only shifting by 1, 2, or 3. */
   ieeel >>= shift;
   ieee2 = (xport2 >> shift) |
   ((xport1 & 0x00000007) << (29 + (3 - shift)));
 /* clear the 1 bit to the left of the binary point */
   ieeel &= 0xffefffff;
 /* set the exponent of the ieee number to be the actual */
   /\!\!^* exponent plus the shift count + 1023. Or this into the ^*/\!\!
   /* first half of the ieee number. The ibm exponent is excess */
   /\!\!^* 64 but is adjusted by 65 since during conversion to ibm */ /\!\!^* format the exponent is incremented by 1 and the fraction */
   /* bits left 4 positions to the right of the radix point. */
   ieee1 |=
   (((((long)(*temp & 0x7f) - 65) << 2) + shift + 1023) << 20)
   (xport1 & 0x80000000);
REVERSE(&ieeel,sizeof(unsigned long));
   memcpy(ieee,((char *)&ieee1)+sizeof(unsigned long)-4,4);
   REVERSE(&ieee2,sizeof(unsigned long));
   memcpy(ieee+4,((char *)&ieee2)+sizeof(unsigned long)-4,4);
   return;
  /* Name: ieee2xpt */
   /* Purpose: converts IEEE to transport */
   /* Usage: rc = ieee2xpt(to_ieee,p_data); */
   /* Notes: this routine is an adaptation of the wzctdbl routine */
   /* from the Apollo. */
                    -----*/
void ieee2xpt(ieee,xport)
unsigned char *ieee; /* ptr to IEEE field (2-8 bytes) */
   unsigned char *xport; /* ptr to xport format (8 bytes) */
register int shift;
   unsigned char misschar;
   int ieee_exp;
   unsigned long xport1, xport2;
  unsigned long ieee1 = 0;
unsigned long ieee2 = 0;
   char ieee8[8];
memcpy(ieee8,ieee,8);
 /*----get 2 longs for shifting-----
   memcpy(((char *)&ieee1)+sizeof(unsigned long)-4,ieee8,4);
   REVERSE & ieee1, sizeof(unsigned long));
memcpy(((char *)&ieee2)+sizeof(unsigned long)-4,ieee8+4,4);
   REVERSE(&ieee2,sizeof(unsigned long));
memset(xport,0,8);
 /*----if IEEE value is missing (1st 2 bytes are FFFF)----*/
   if (*ieee8 == (char)0xff && ieee8[1] == (char)0xff) {
   misschar = ~ieee8[2];
   *xport = (misschar == 0xD2) ? 0x6D : misschar;
   return;
```

```
/* Translate IEEE floating point number into IBM format float */
  /* IEEE format: */
  /* 6 5 0 */
  /* 3 1 0 */
/* */
  /* SEEEEEEEEEEEMMMM ..... MMMM */
  /\!\!\,^* Sign bit, 11 bit exponent, 52 fraction. Exponent is excess ^*/
  /* 1023. The fraction is multiplied by a power of 2 of the */
  /st actual exponent. Normalized floating point numbers are st/
  /* represented with the binary point immediately to the left */
  /* of the fraction with an implied "1" to the left of the */
  /* binary point. */
  /* IBM format: */
  /* 6 5 0 */
  /* 3 5 0 */
  /* */
  /* SEEEEEEEMMMM ..... MMMM */
/* */
  /* Sign bit, 7 bit exponent, 56 bit fraction. Exponent is */ /* excess 64. The fraction is multiplied by a power of 16 of */
  /* of the actual exponent. Normalized floating point numbers */
  /* are presented with the radix point immediately to the left */
  /* of the high order hex fraction digit. */
  /* How do you translate from local to IBM format? */
  /* The ieee format gives you a number that has a power of 2 */ \,
  /* exponent and a fraction of the form "1.<fraction bits>". */
  /* The first step is to get that "1" bit back into the */
  /* fraction. Right shift it down 1 position, set the high */
  /* order bit and reduce the binary exponent by 1. Now we have */ /* a fraction that looks like ".1<fraction bits>" and it's */ /* ready to be shoved into ibm format. The ibm fraction has 4 */
  /* more bits than the ieee, the ieee fraction must therefore
  /* be shifted left 4 positions before moving it in. We must */
  /* also correct the fraction bits to account for the loss of 2*/
  /* bits when converting from a binary exponent to a hex one */ /* (>> 2). We must shift the fraction left for 0, 1, 2, or 3 */
  /* positions to maintain the proper magnitude. Doing */
  /* conversion this way would tend to lose bits in the fraction*/
  /* which is not desirable or necessary if we cheat a bit. */
  /st First of all, we know that we are going to have to shift st/
  /* the ieee fraction left 4 places to put it in the right */
  /* position; we won't do that, we'll just leave it where it is*/
/* and increment the ibm exponent by one, this will have the */
  /* same effect and we won't have to do any shifting. Now, *,
  /* since we have 4 bits in front of the fraction to work with,*/
  /^{\star} we won't lose any bits. We set the bit to the left of the ^{\star}/ /^{\star} fraction which is the implicit "1" in the ieee fraction. We*/
  /* then adjust the fraction to account for the loss of bits */
  /* when going to a hex exponent. This adjustment will never *
  /* involve shifting by more than 3 positions so no bits are */
  /* lost. */
/\!\!\!\!\!\!^* Get ieee number less the exponent into the first half of ^*/\!\!\!\!
  /* the ibm number */
xport1 = ieee1 & 0x000fffff;
/st get the second half of the number into the second half of st/
     the ibm number and see if both halves are 0. If so, ibm is */
  /* also 0 and we just return */
if ((!(xport2 = ieee2)) && !ieee1) {
  ieee_exp = 0;
  goto doret;
/* get the actual exponent value out of the ieee number. The */
     ibm fraction is a power of 16 and the ieee fraction a power*/
  /* of 2 (16 ** n == 2 ** 4n). Save the low order 2 bits since */
  /* they will get lost when we divide the exponent by 4 (right */
  /* shift by 2) and we will have to shift the fraction by the */ /* appropriate number of bits to keep the proper magnitude. */
  (ieee\_exp = (int)(((ieee1 >> 16) \& 0x7ff0) >> 4) - 1023)
```

```
/* the ieee format has an implied "1" immdeiately to the left */
   /* of the binary point. Show it in here. */
xport1 |= 0x00100000;
 if (shift)
   {
/* set the first half of the ibm number by shifting it left*/
   /* the appropriate number of bits and oring in the bits */
   /* from the lower half that would have been shifted in (if */
   /* we could shift a double). The shift count can never */
   /* exceed 3, so all we care about are the high order 3 */
/* bits. We don't want sign extention so make sure it's an */
   /* unsigned char. We'll shift either5, 6, or 7 places to */
   /* keep 3, 2, or 1 bits. After that, shift the second half */
   /st of the number the right number of places. We always get st/
   /* zero fill on left shifts. */
   xport1 = (xport1 << shift) |
((unsigned char) (((ieee2 >> 24) & 0xE0) >>
   (5 + (3 - shift)));
 xport2 <<= shift;</pre>
 /\,^{\star} Now set the ibm exponent and the sign of the fraction. The ^{\star}/\,
   /^{\star} power of 2 ieee exponent must be divided by 4 and made ^{\star}/ /^{\star} excess 64 (we add 65 here because of the poisition of the ^{\star}/
   /* fraction bits, essentially 4 positions lower than they */
   /* should be so we incrment the ibm exponent). */
   (((ieee_exp >>2) + 65) | ((ieee1 >> 24) & 0x80)) << 24;
 /* If the ieee exponent is greater than 248 or less than -260,*/
   /* then it cannot fit in the ibm exponent field. Send back the*/
   /* appropriate flag. */
   if (-260 <= ieee_exp && ieee_exp <= 248) {
   REVERSE(&xport1,sizeof(unsigned long));
   memcpy(xport,((char *)&xport1)+sizeof(unsigned long)-4,4);
   REVERSE(&xport2, sizeof(unsigned long));
   memcpy(xport+4,((char *)&xport2)+sizeof(unsigned long)-4,4);
   memset(xport,0xFF,8);
 if (ieee_exp > 248)
   *xport = 0x7f;
   return;
#ifdef DEFINE REVERSE
   void REVERSE(intp,1)
   char *intp;
   int 1;
   int i,j;
   char save;
   static int one = 1;
#if !defined(BIG_ENDIAN) && !defined(LITTLE_ENDIAN)
   if (((unsigned char *)&one)[sizeof(one)-1] == 1)
   return;
   #endif
j = 1/2;
  for (i=0;i<j;i++) {
   save = intp[i];
   intp[i] = intp[l-i-1];
   intp[1-i-1] = save;
   #endif
```